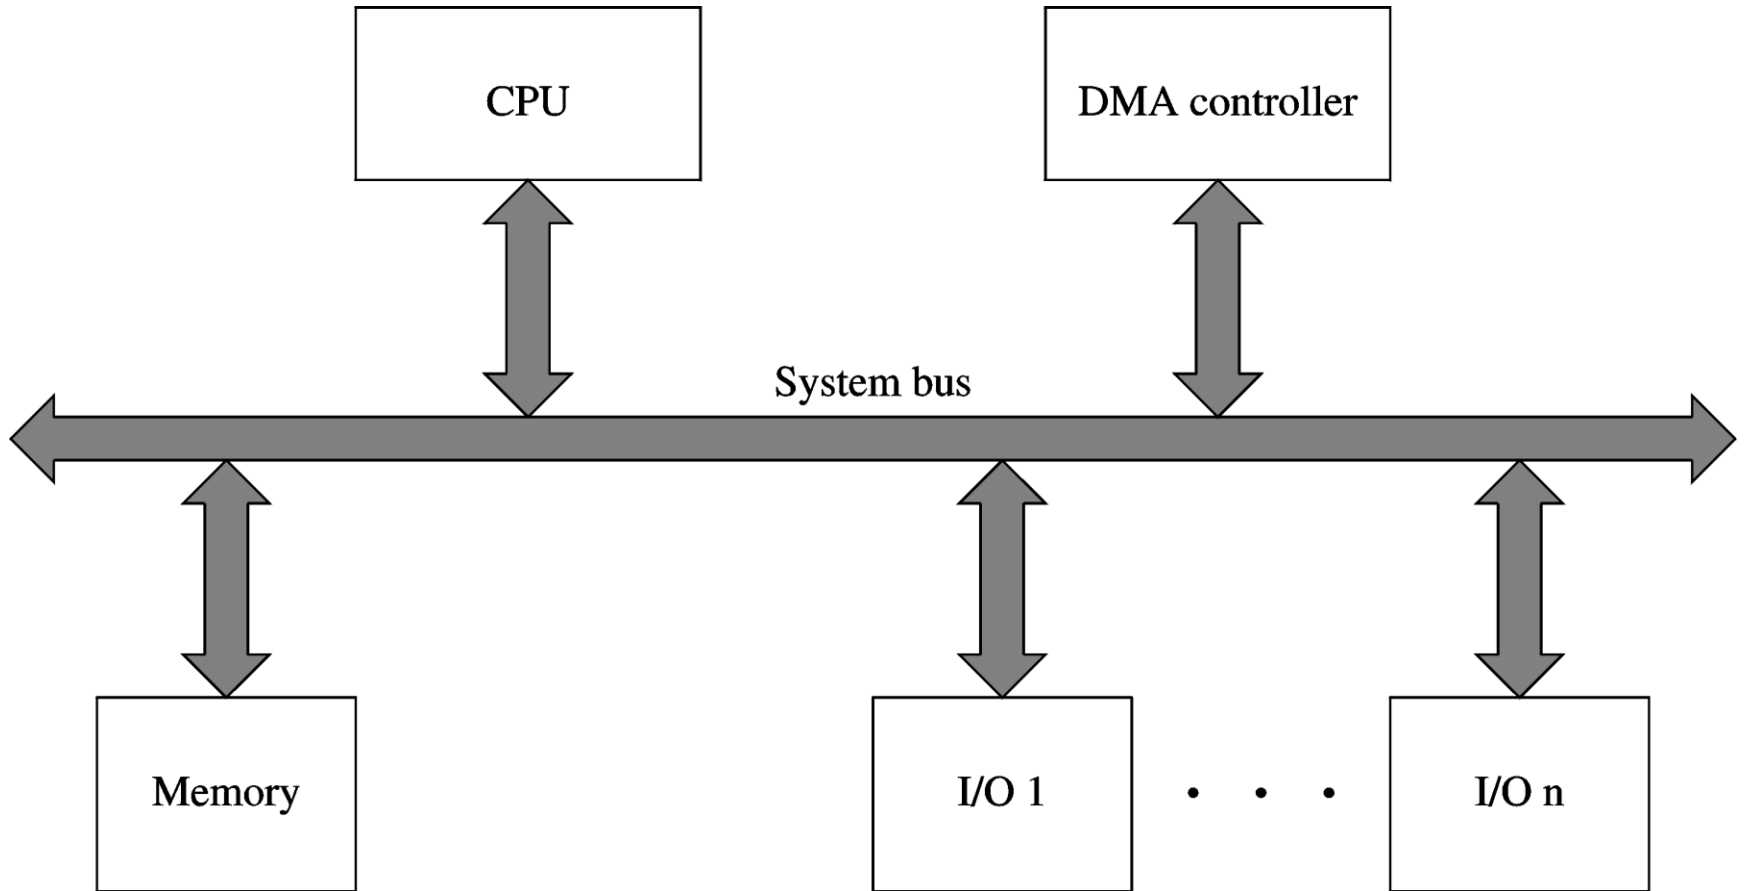


System Bus

Outline

- Introduction
- Bus design issues
 - Bus width
 - Bus type
 - Bus operations
- Synchronous bus
 - Bus operation
 - Wait states
 - Block transfer
- Asynchronous bus
- Bus arbitration
 - Dynamic bus arbitration
 - Implementation
 - Centralized
 - Distributed
- Example buses
 - ISA
 - PCI
 - AGP
 - PCI-X
 - PCMCIA

Introduction



- System bus consists of
 - Address bus
 - Data bus
 - Control bus
- Buses can be
 - Dedicated
 - Multiplexed
 - Synchronous
 - Asynchronous

Introduction (cont'd)

- Control bus
 - Memory read and Memory write
 - I/O read and I/O write
 - Ready
 - Bus request and Bus grant
 - Interrupt and Interrupt acknowledgement
 - DMA request and DMA acknowledgement
 - Clock
 - Reset

Bus Design Issues

- Need to consider several design issues
 - Bus width
 - Data and address buses
 - Bus type
 - Dedicated or multiplexed
 - Bus operations
 - Read, write, block transfer, interrupt, ...
 - Bus arbitration
 - Centralized or distributed
 - Bus timing
 - Synchronous or asynchronous

- **Data bus width**

- A critical parameter in determining system performance
- Need not correspond to the ISA-specific value
 - Pentium is a 32-bit processor
 - But has 64-bit data bus
 - Itanium is a 64-bit processor
 - But has 128-bit data bus
- The wider the data bus, the better
 - Wider buses are expensive

- **Address bus width**

- Determines the system addressing capacity
- N address lines directly address 2^N memory locations

- 8086: 20 address lines
 - Could address 1 MB of memory
- Pentium: 32 address lines
 - Could address 4 GB of memory
- Itanium: 64 address lines
 - Could address 2^{64} bytes of memory

Bus Type

- **Dedicated buses**
 - Separate buses dedicated to carry data and address information
 - Good for performance
 - But increases cost
- **Multiplexed buses**
 - Data and address information is time multiplexed on a shared bus
 - Better utilization of buses
 - Reduces cost

Bus Operations

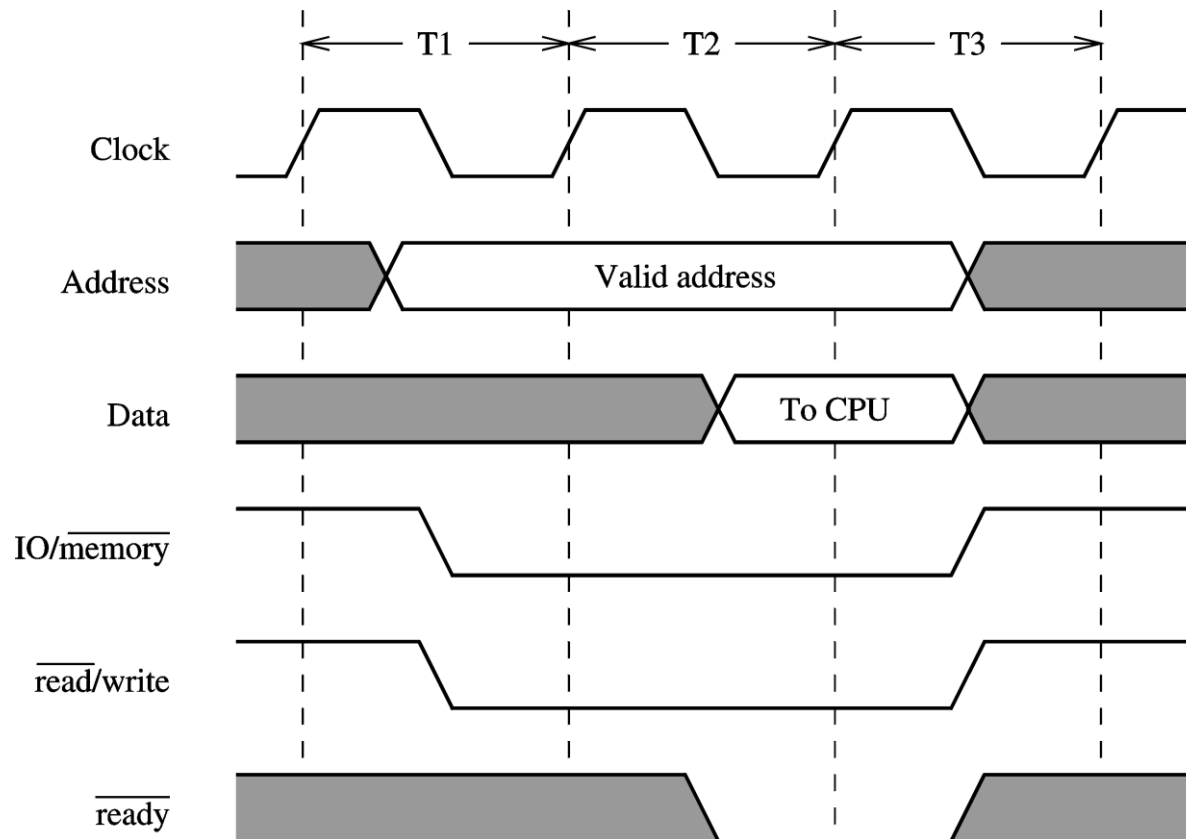
- Basic operations
 - Read and write
- Block transfer operations
 - Read or write several contiguous memory locations
 - Example: cache line fill
- Read-modify-write operation
 - Useful for critical sections
- Interrupt operation
- Several other types...

Synchronous Bus

- A bus clock signal provides timing information for all actions
 - Changes occur relative to the falling or rising edge of the clock
 - Choosing appropriate clock is important
 - Easier to implement
 - Most buses are synchronous
- Bus operations can be
 - With no wait states, or
 - With wait states


Synchronous Bus (cont'd)

- Memory read operation with no wait states



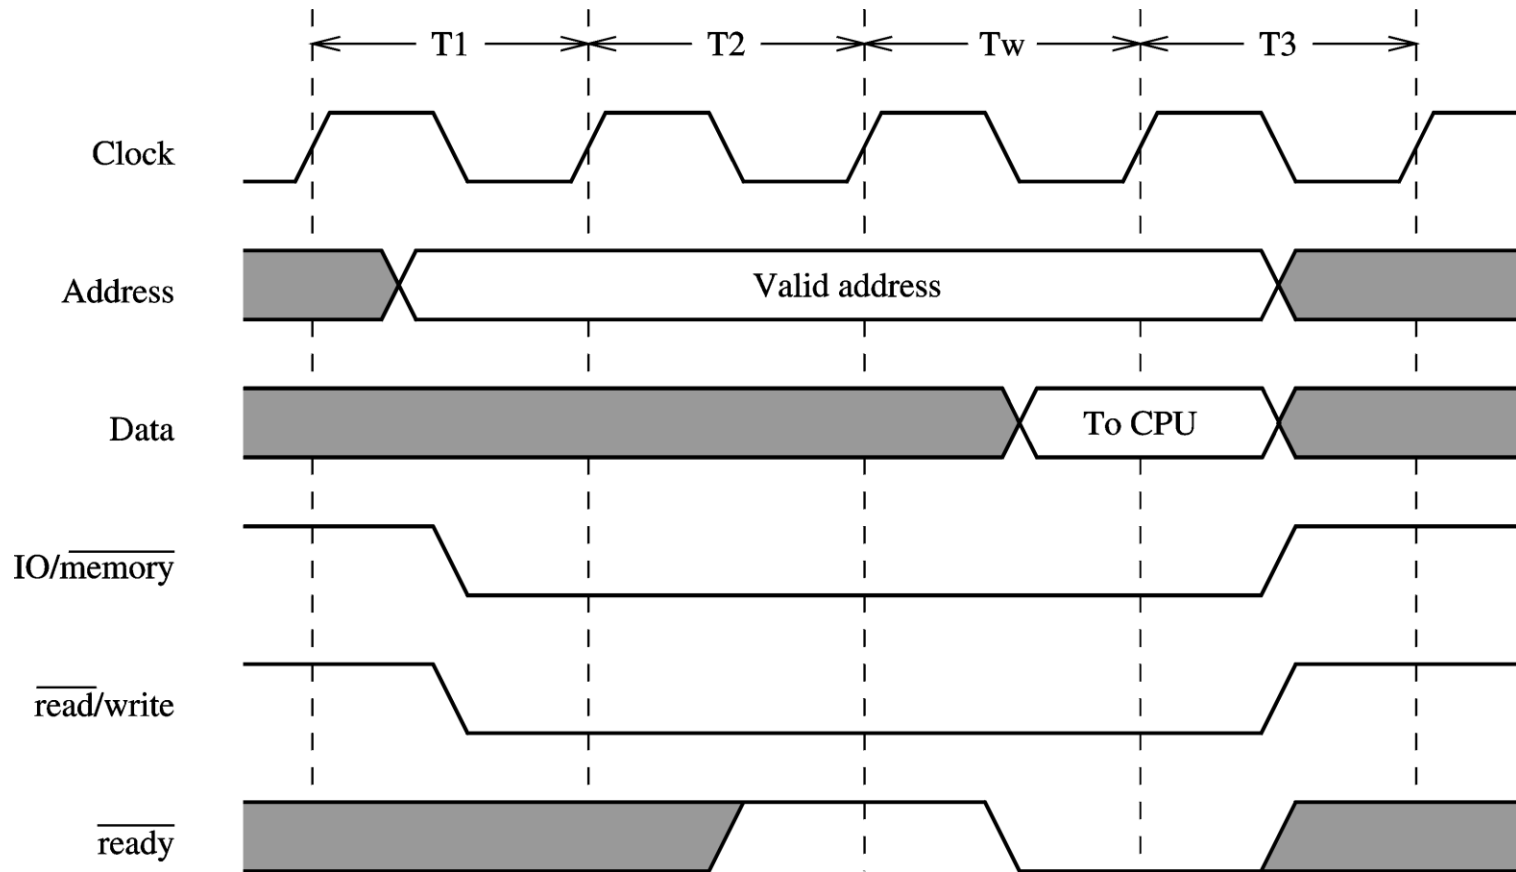
Synchronous Bus (cont'd)

- Memory write operation with no wait states



Synchronous Bus (cont'd)

- Memory read operation with a wait state

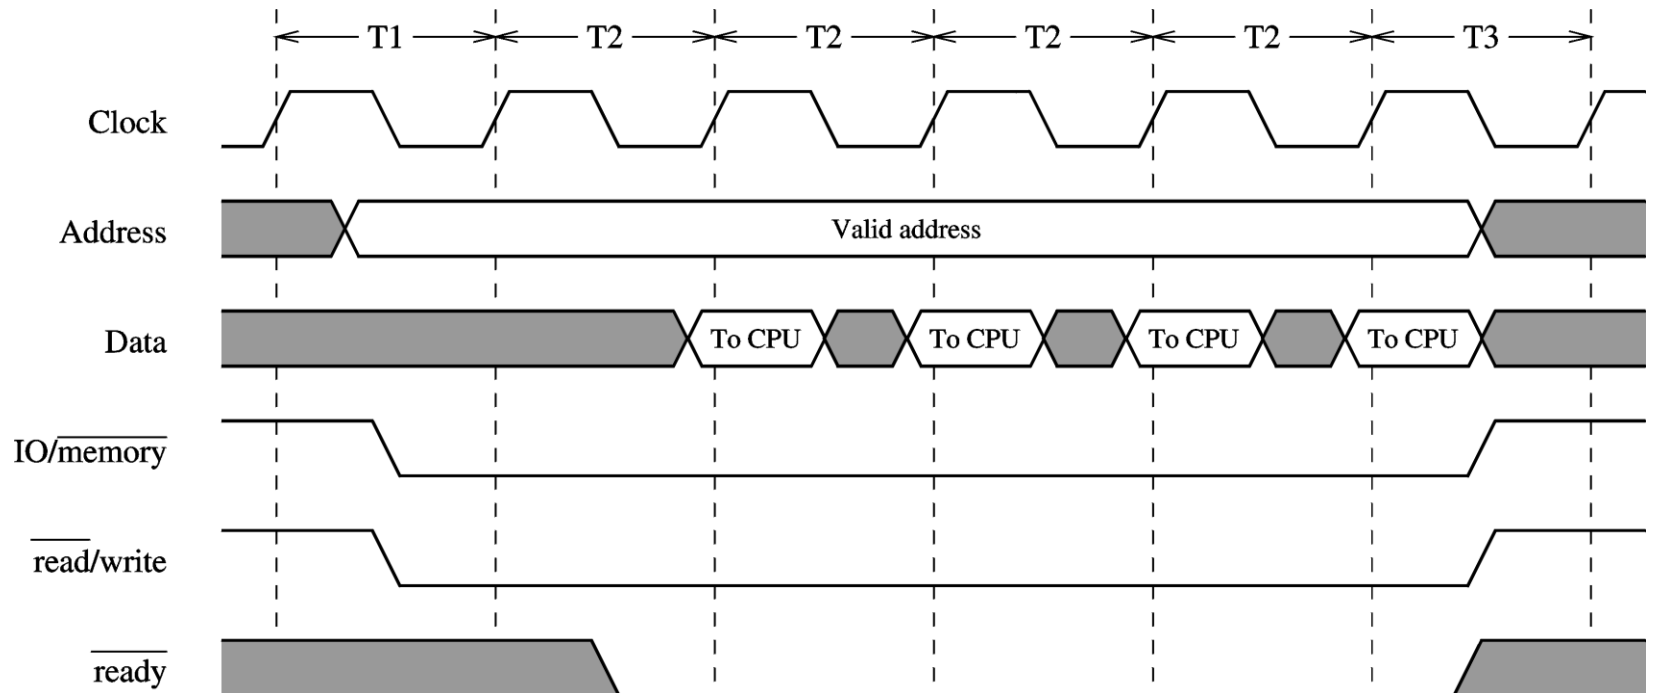


Synchronous Bus (cont'd)

- Memory write operation with a wait state

Synchronous Bus (cont'd)

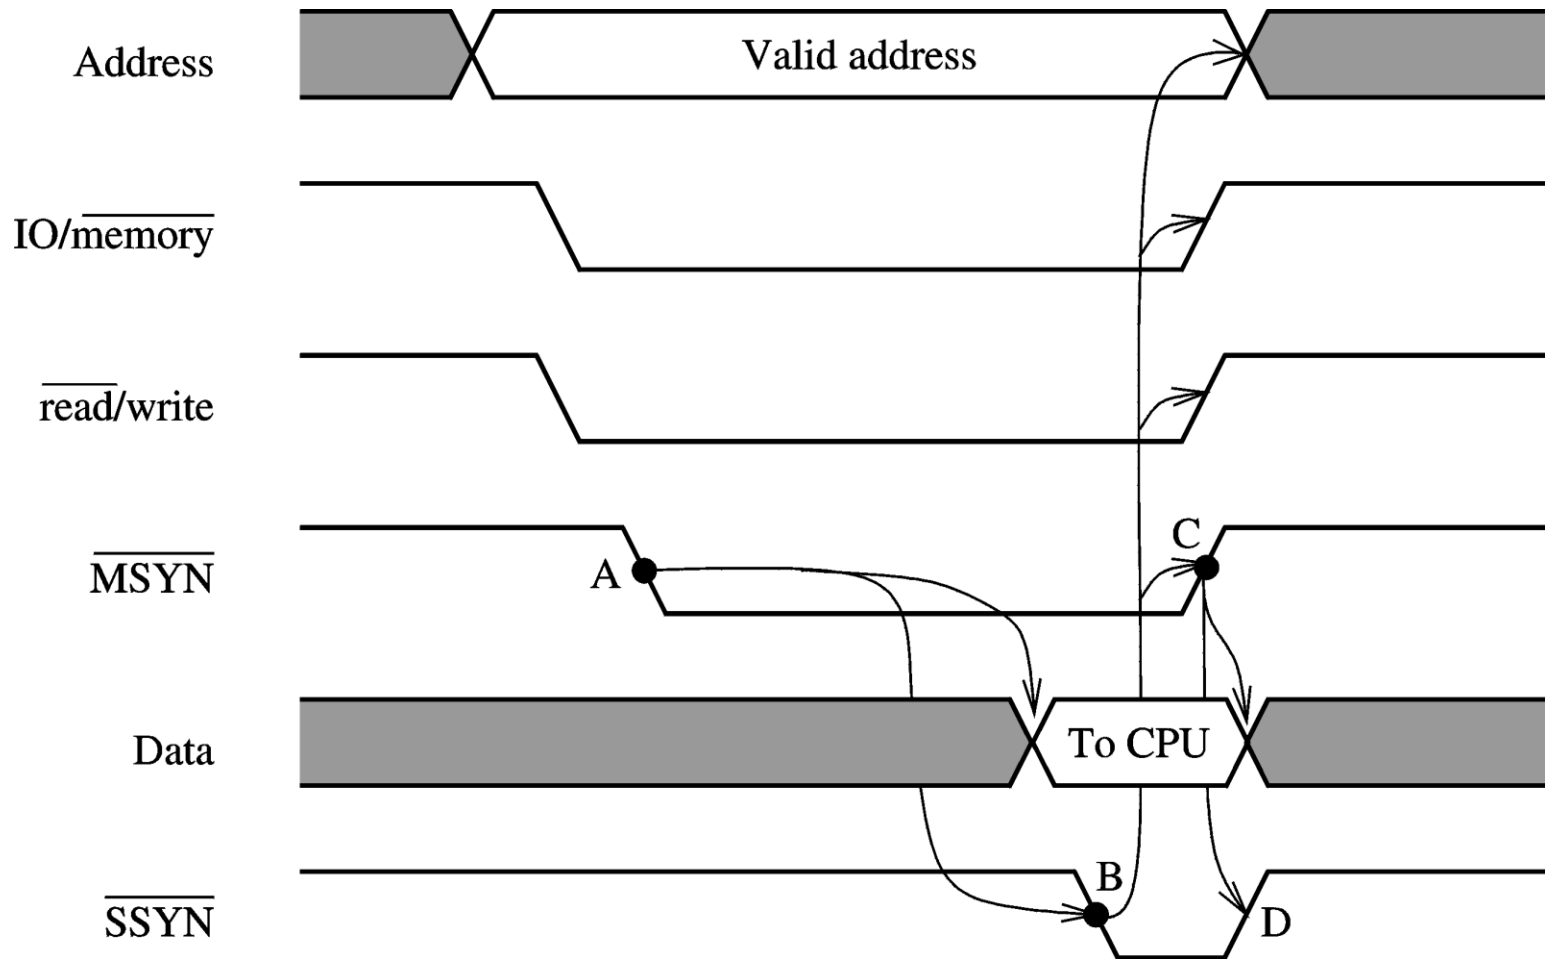
- Block transfer of data



Asynchronous Bus

- No clock signal to synchronize actions
- Operates in master-slave mode
- Uses handshaking to perform a bus transaction
 - Two synchronization signals facilitate this
 - Master synchronization (MSYN)
 - Slave synchronization (SSYN)
- Advantage of asynchronous buses
 - No need for bus clock
- Synchronous buses
 - Easier to implement

Asynchronous Bus (cont'd)



Bus Arbitration

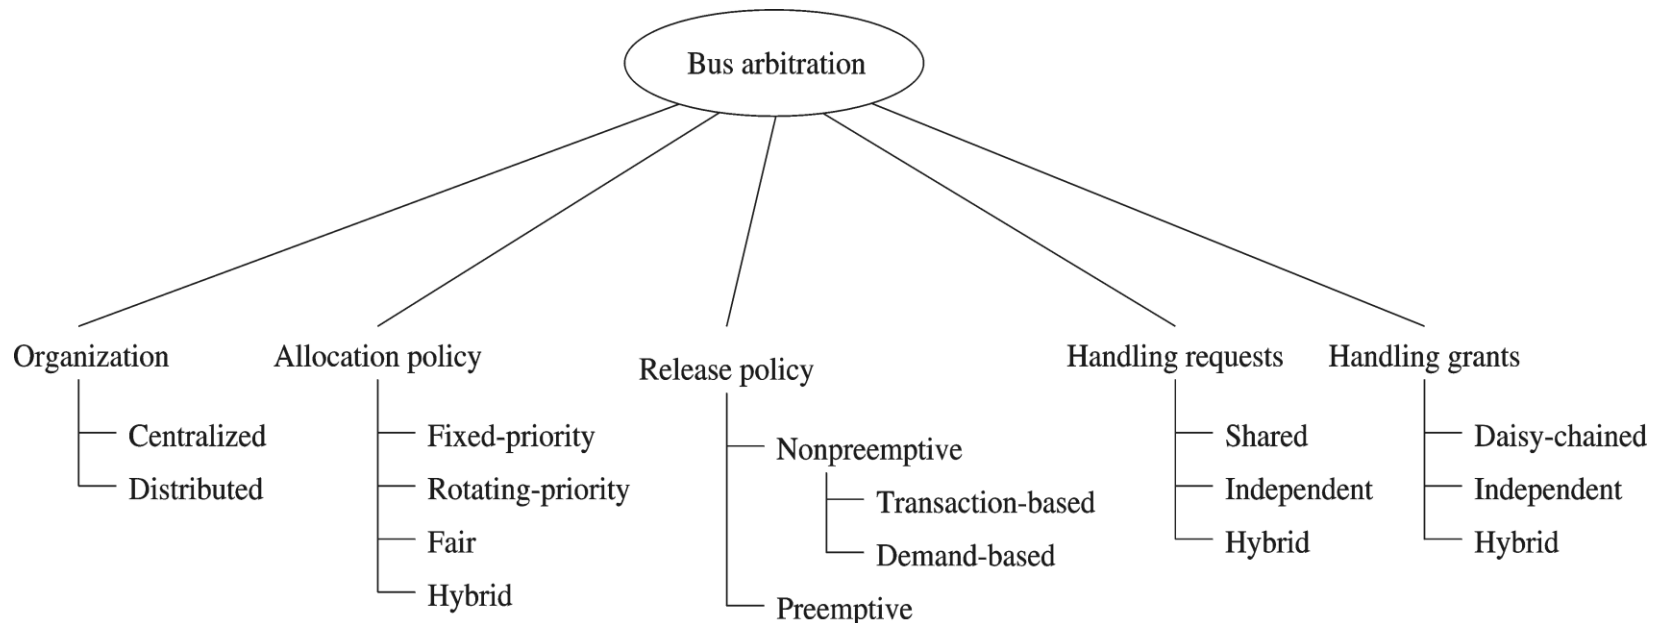
- More than one bus master can request the bus
 - Need an arbitration mechanism to allocate the bus
- Bus arbitration can be done either
 - Statically
 - Dynamically
- Static arbitration
 - Done in a predetermined way
 - Easy to implement
 - Does not take needs into account
 - Poor utilization
 - Bus could be assigned even when not needed

Dynamic Bus Arbitration

- Bus allocated only in response to a request
- Each master is equipped with
 - Bus request line
 - Bus grant line
- A master uses the bus request line to let others know that it needs the bus
- Before a master can use the bus, it must receive permission to use the bus via the bus grant line

- Bus arbitration can be implemented
 - Centralized
 - Distributed
- Centralized arbitration
 - A central arbiter receives all bus requests
 - Uses an allocation policy to determine which request should be granted
 - This decision is conveyed through the bus grant lines
 - Once the transaction is over, bus is released
 - A bus release policy determines the actual release mechanism

- Distributed arbitration
 - Arbitration hardware is distributed among the masters
 - A distributed arbitration algorithm is used to determine who should get the bus



- Bus Allocation Policies
 - Fixed priority
 - Each master is assigned a fixed priority
 - Highest priority master always gets the bus
 - Priorities can be assigned based on the importance of service
 - Rotating priority
 - Priority is not fixed
 - Several ways of changing priority
 - Increase the priority as a function of waiting time
 - Lowest priority for the master that just received the bus

- Bus Allocation Policies

- Fair policies

- A fair policy will not allow starvation
 - Rotating priority policies are fair
 - Fair policies need not use priorities
 - Fairness can be defined in several ways
 - A window-based request satisfaction
 - Within a specified time period
 - » In PCI, we can specify the maximum delay to grant a request

- Hybrid policies

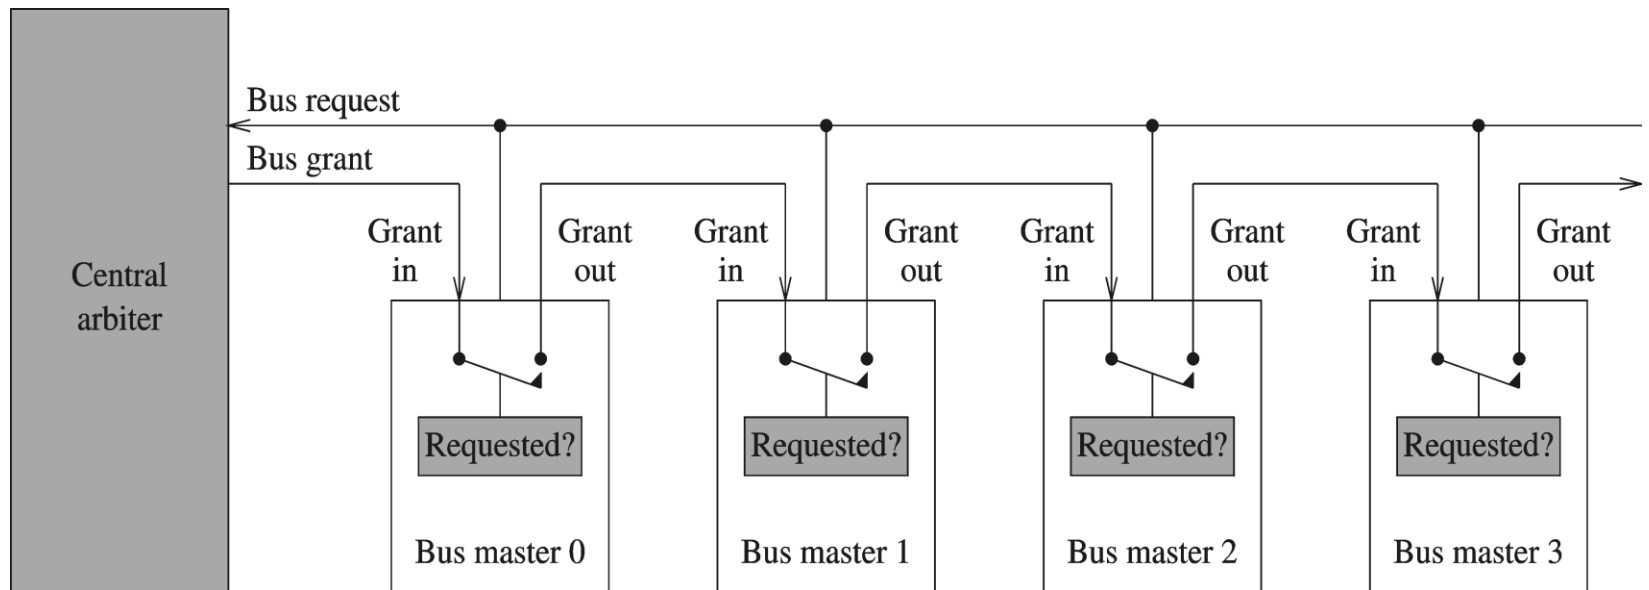
- Both priority and fairness can be incorporated into a single policy

- Bus Release Policies
 - Governs the conditions under which the current master releases the bus
 - Two types
 - Non-preemptive
 - Current master voluntarily releases the bus
 - Disadvantage
 - » May hold bus for long time
 - Preemptive
 - Forces the current master to release the bus without completing its bus transaction

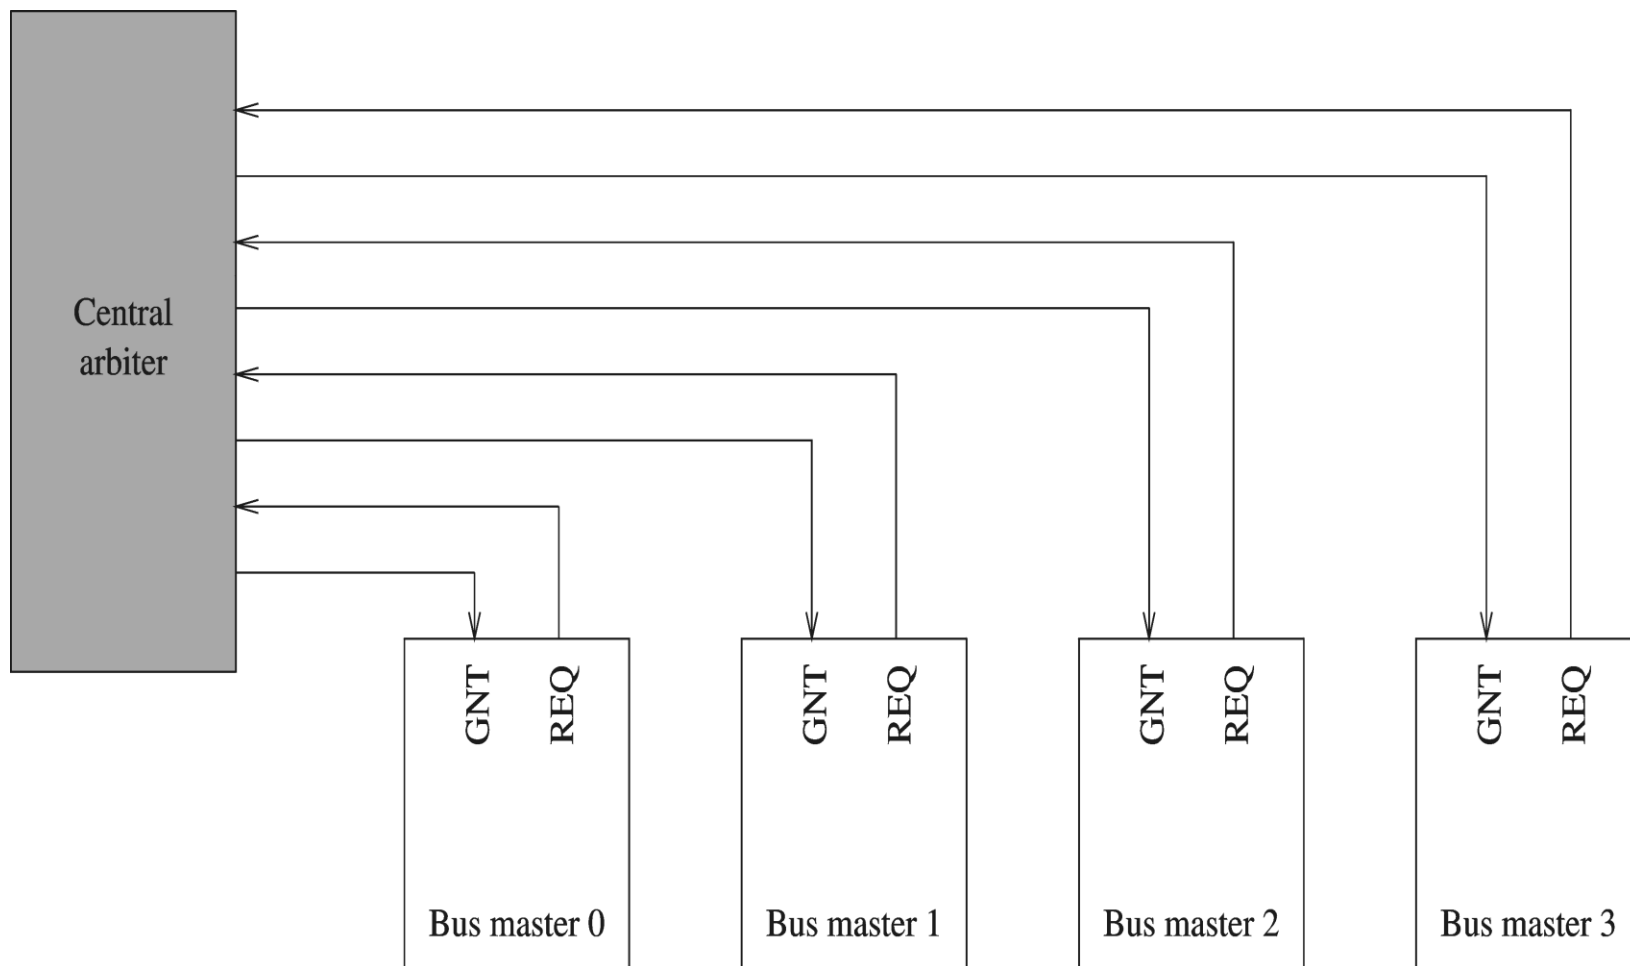
- Non-Preemptive Bus Release Policies
 - Transaction-based release
 - Releases bus after completing the current transaction
 - Requests bus again if it has more transactions
 - By releasing the bus after every transactions, fairness can be ensured
 - Easy to implement
 - Unnecessary overhead if only one master needs the bus
 - Demand-driven release
 - Avoids unnecessary bus requests of the previous policy
 - Releases the bus only if another master requests the bus
 - More efficient

- Implementation
 - Centralized bus arbitration
 - Daisy-chaining
 - Uses a single, shared bus request signal
 - Central arbiter sends the grant signal to the first master in the chain
 - » Each master passes the grant signal to its neighbor if it does need the bus
 - » Grabs the grant signal if it wants the bus
 - Easy to implement
 - » Needs three control lines independent of the number of hosts

- Three problems
 - Implements fixed priority
 - Arbitration time proportional to number of hosts
 - Not fault-tolerant




- Independent requests
 - Arbiter is connected to each master
 - Variety of bus allocation policies can be implemented
 - Avoids the pitfalls of daisy-chaining
 - Fault-tolerant
 - Short, constant arbitration time
 - Disadvantages
 - Complex to implement
 - Number of control signals is proportional to the number of hosts
 - PCI uses this implementation technique



- Hybrid scheme
 - Daisy-chaining and independent requests represent two extremes
 - Daisy-chaining
 - Cheaper but several problems
 - Independent requests
 - Expensive but avoids the problems of daisy-chaining
 - Hybrid scheme combines good features of these two
 - Bus masters are divided into N classes
 - Independent strategy at the class-level
 - Daisy-chaining within each class
 - VME bus uses a hybrid policy

Dynamic Bus Arbitration (cont'd)

VME bus arbitration



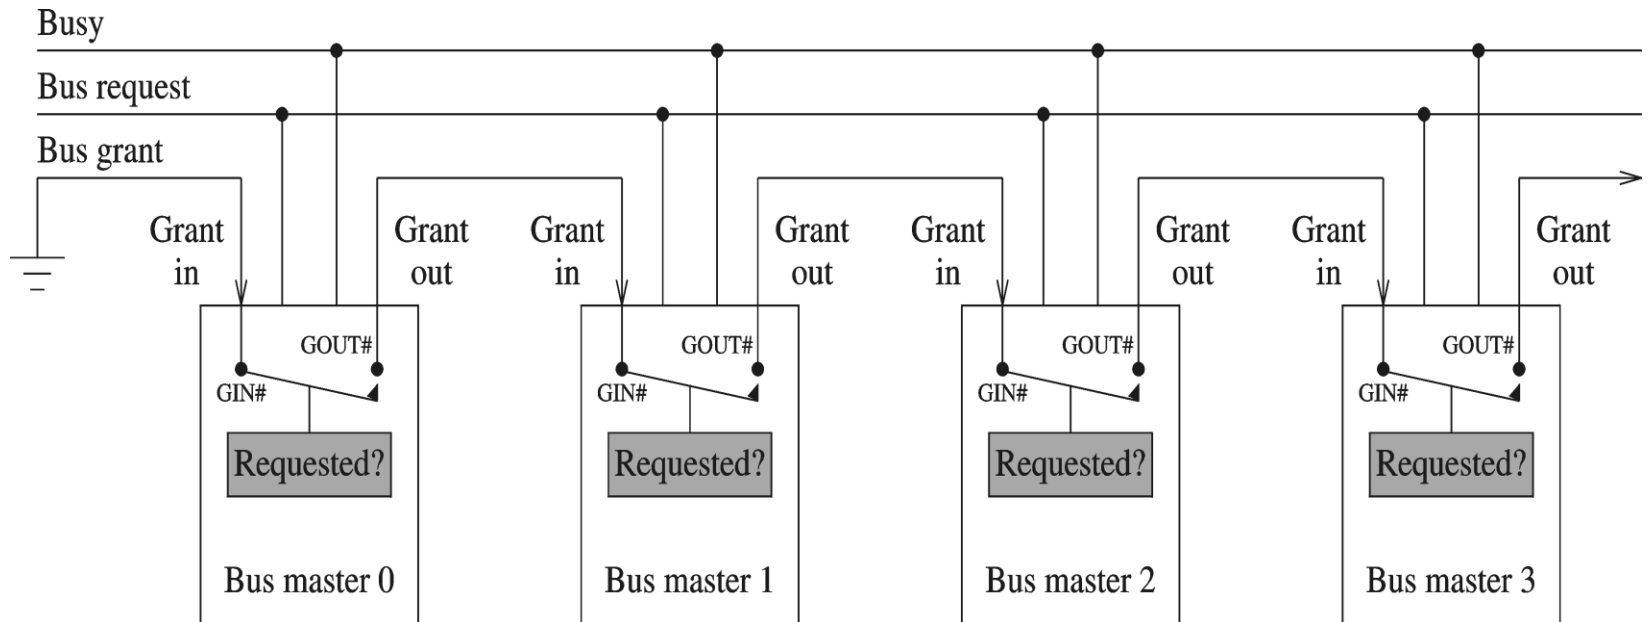
Dynamic Bus Arbitration (cont'd)

- VME bus arbitration
 - Allocation policies supported
 - Fixed-priority
 - GR0: Lowest priority
 - GR3: Highest priority
 - Rotating-priority
 - Round-robin based policy
 - » Assigns the lowest priority to the bus that just received the bus
 - Daisy-chaining
 - Implemented by connecting all masters to BR3 grant request line

Dynamic Bus Arbitration (cont'd)

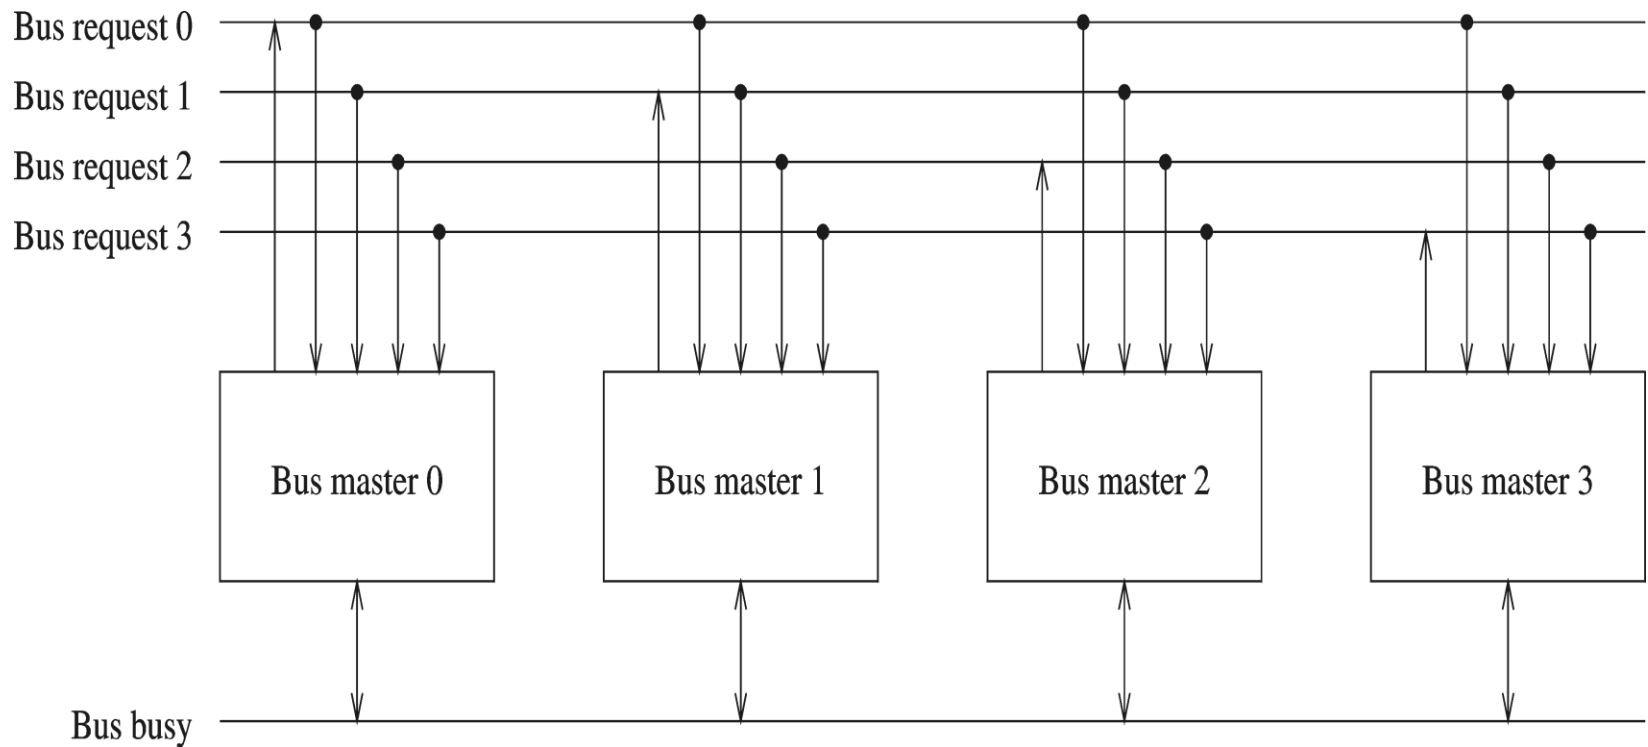
- VME bus arbitration (cont'd)
 - Release policies supported
 - Default release policy
 - Transaction-based
 - Non-preemptive
 - When fixed priority is used
 - Preemptive release is possible when a higher priority request comes in
 - To effect preemption
 - » Bus clear (BCLR) line is asserted
 - » The current bus master releases the bus when BCLR is low

- Distributed bus arbitration
 - Daisy-chaining
 - Needs three control lines as before



— Separate bus request lines

- All bus masters can read all bus request lines
- Highest priority master gets the bus



- Separate bus request lines (cont'd)
 - Implements fixed-priority scheme
 - Starvation is a potential problem
 - To avoid starvation, a kind of honor system can be used
 - The highest-priority master that just used the bus will not raise its bus request until all the other lower priority masters have been allocated the bus

Example Buses


- We look at five buses
 - ISA
 - Supports older text-based applications
 - PCI
 - Supports modern window-based systems
 - AGP
 - Supports high-performance graphics and full-motion video
 - PCI-X
 - Improved and faster PCI
 - PCMCIA
 - Useful for laptops

ISA Bus

- Closely associated with the PC system bus
 - ISA = Industry Standard Architecture
- First ISA bus (8-bit wide data path)
 - Based on 8088 processor
 - It had 82 pins including
 - 20 address lines
 - 8 data lines
 - 6 interrupt signals
 - Memory read, memory write
 - I/O read, and I/O write
 - 4 DMA requests and 4 DMA acks

ISA Bus (cont'd)

- 16-bit ISA
 - Added 36 pins to the 8-bit ISA bus
 - 24 address lines
 - 16 data lines
 - Backward compatible with 8-bit ISA

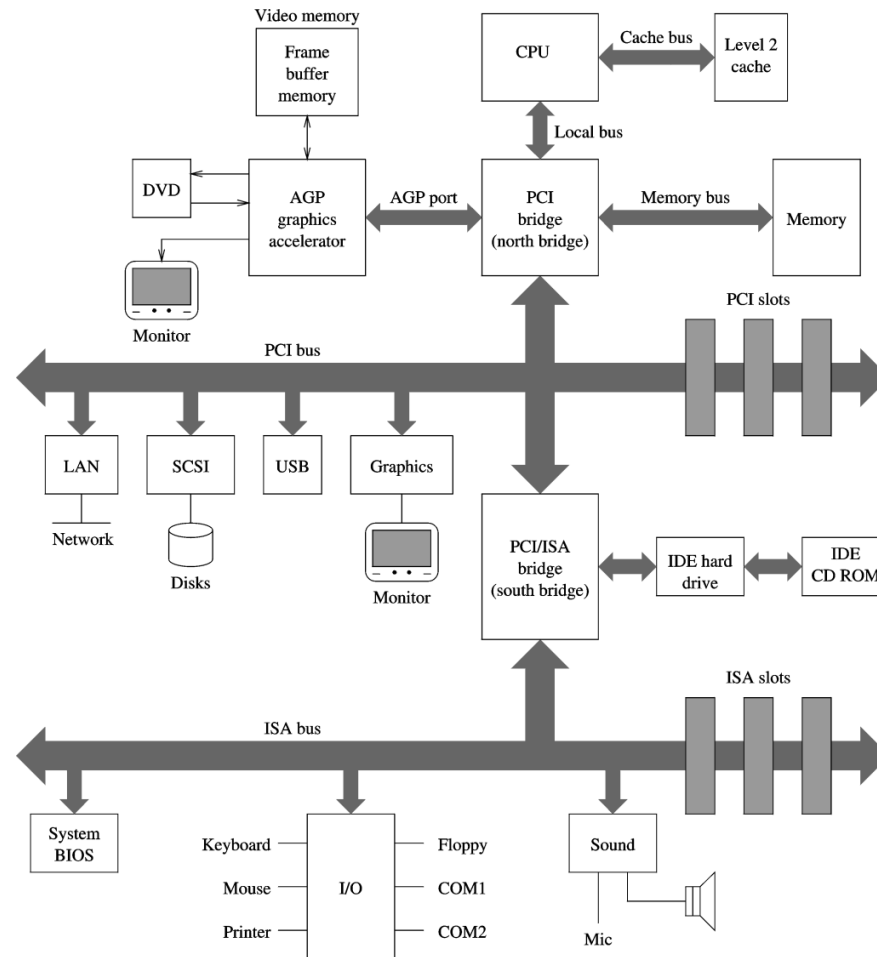


ISA Bus (cont'd)

- Operates at 8.33 MHz
- Bandwidth of about 8 MB/s
- 32-bit processors need more support
 - Several attempts were made to accommodate
 - EISA (Extended ISA)
 - Bus mastering signals
 - MCA (Micro Channel Architecture)
 - IBM proprietary
 - Never really took off
- ISA is used for older, slower I/O devices

PC System Buses

ISA
PCI
AGP



PCI Bus



- Work began in 1990
 - Intel began work on a new bus for their Pentium systems
 - Processor independent
 - To support higher bandwidth requirements of window-based systems
 - Original version (1.0) developed by Intel in 1990
 - Version 2 in 1993
 - Version 2.1 in 1995
 - Version 2.2 introduced power management for mobile systems

PCI Bus (cont'd)

- Implemented either
 - 32-bit or 64-bit
- Operate at
 - 33 MHz or 66 MHz
 - 5 V (older cards) or 3.3 V (newer ones)
- 32-bit PCI operating at 33 MHz
 - Provides peak bandwidth of 133 MB/s
- 64-bit PCI operating at 66 MHz
 - Provides peak bandwidth of 528 MB/s



PCI Bus (cont'd)

3.3 V 64bit connector



3.3 V 32bit connector

5 V 32bit connector



5 V 64bit connector

PCI Bus (cont'd)

- Bus signals
 - Mandatory signals
 - System signals
 - Transaction control signals
 - Bus arbitration and error reporting signals
 - Optional signals
 - 64 bit extension signals and Interrupt request signals
- System signals
 - Clock (CLK)
 - Reset (RST#)
 - Address/Data bus (AD[0-31])
 - 32 lines multiplexed to serve both address and data

PCI Bus (cont'd)

- Command Bus (C/BE#[0-3])
 - Multiplexed command and byte enable (BE#) signals
 - Command signals identify transaction type
 - » Memory read, memory write, ...
 - » Asserted during address phase
 - BE signals select the bytes to be transferred
 - » Asserted during data phase
 - » BE0# identifies byte 0, BE#1 byte 1, ...
 - » 0000 = all four bytes
 - » 1111 = none of the bytes (null data phase)
- Parity Signal (PAR)
 - Even parity for AD and C/BE lines

PCI Bus (cont'd)

- Transaction Control Signals
 - Cycle Frame (FRAME#)
 - Indicates start of a bus transaction
 - Also indicates the length of the bus transaction cycle
 - Initiator Ready (IRDY#)
 - During Write transaction:
 - Indicates the initiator has placed data on AD lines
 - During Read transaction:
 - Indicates the initiator is ready to accept data
 - Target Ready (TRDY#)
 - Complements IRDY# signal
 - IRDY# and TRDY# together implement handshake to transfer data

PCI Bus (cont'd)

- Transaction Control Signals
 - Stop Transaction (STOP#)
 - Target asserts this to tell initiator that it wants to terminate the current transaction
 - Initialization Device Select (IDSEL)
 - Used as a chip select for configuration read and write transactions
 - Device Select (DEVSEL#)
 - Selected target asserts this to tell the initiator that it is present
 - Bus Lock (LOCK#)
 - Initiator uses this to lock the target to execute atomic

PCI Bus (cont'd)

- Bus Arbitration Signals
 - Uses centralized bus arbitration with independent request and grant lines
 - Bus Request (REQ#)
 - A device asserts when it needs the bus
 - Bus Grant (GNT#)
 - Bus arbiter asserts this signal to indicate allocation of the bus
 - Bus arbitration can overlap execution of another transaction
 - Improves PCI performance

PCI Bus (cont'd)

- Error Reporting Signals
 - Parity Error (PERR#)
 - All devices are expected to report this error
 - Exceptions exist
 - E.g. when transmitting video frames
 - System Error (SERR#)
 - Any PCI device can generate this signal
 - To indicate address and other errors
 - Typically connected to NMI (non-maskable interrupt)

PCI Bus (cont'd)

- 64-bit Extension Signals
 - Address/Data Lines (AD[32 - 63])
 - Extension to 64 bits
 - Command bus (C/BE#[4 - 7])
 - Extended by 4 lines to handle 8 bytes
 - Request 64-Bit Transfer (REQ64#)
 - Indicates to target that initiator likes 64-bit transfers
 - Acknowledge 64-Bit Transfer (ACK64#)
 - Target indicates that it is capable of 64-bit transfers
 - Parity Bit for Upper Data (PAR64)
 - Even parity for upper 32 AD bits and four command

PCI Bus (cont'd)

- Interrupt Request Lines
 - Four interrupt lines
 - INTA#, INTB#, INTC#, INTD#
 - Not shared
- Additional signals
 - To support snoopy cache protocol
 - IEEE 1149.1 boundary scan signals
 - Allows in-circuit testing of PCI devices
 - M66En signal to indicate bus frequency
 - Low: 33 MHz
 - High: 66 MHz

PCI Commands

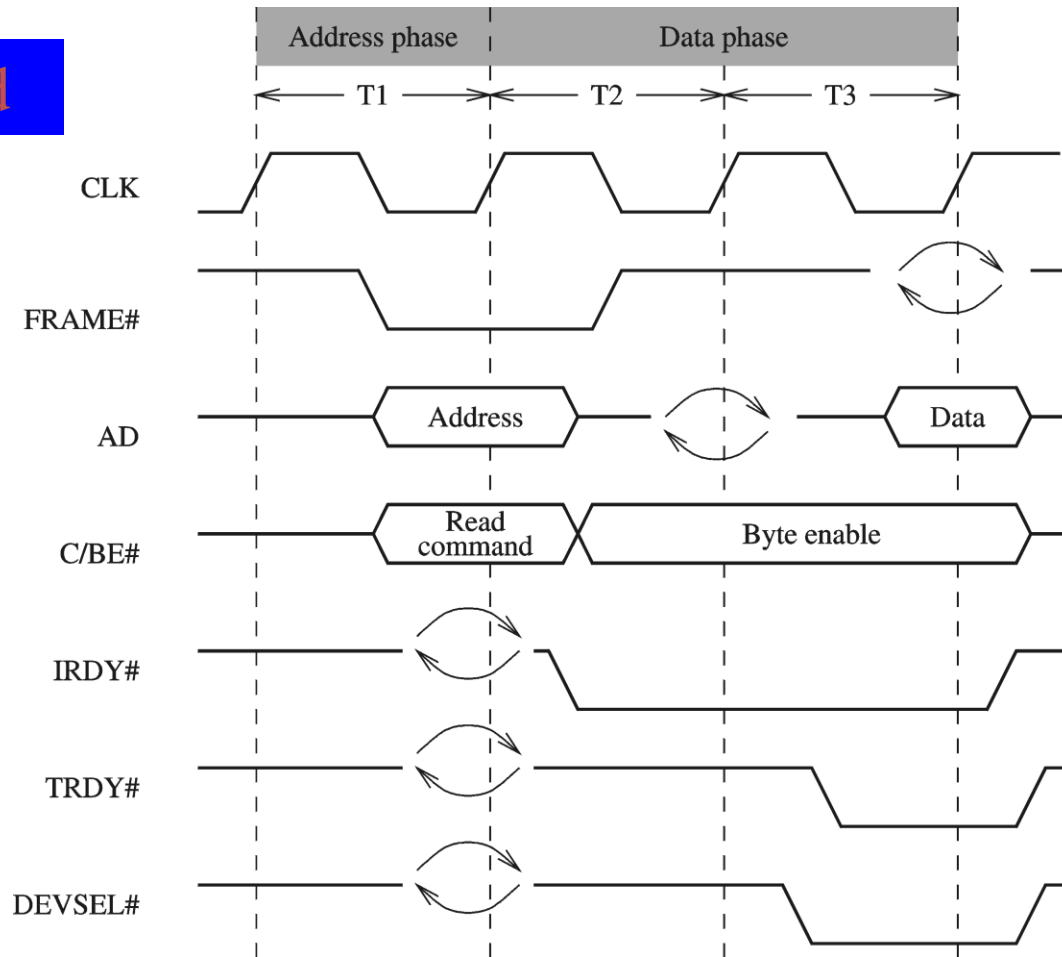
- A command value is placed on C/BE lines during the address phase
 - I/O operations
 - I/O Read and I/O Write
 - Memory operations
 - Standard memory operations
 - Memory Read and Memory Write
 - Bulk memory operations
 - Memory Read Line
 - Memory Read Multiple
 - Memory Write-and-Invalidate

PCI Commands (cont'd)

- Configuration operations
 - Every PCI device has a 256-byte configuration space
 - Two commands:
 - Configuration Read and Configuration Write
- Miscellaneous operations
 - Special Cycle Command
 - Used to broadcast a message to all PCI targets
 - Shutdown and Halt
 - Dual Address Cycle Command
 - Allows 32-bit initiator to use 64-bit addresses

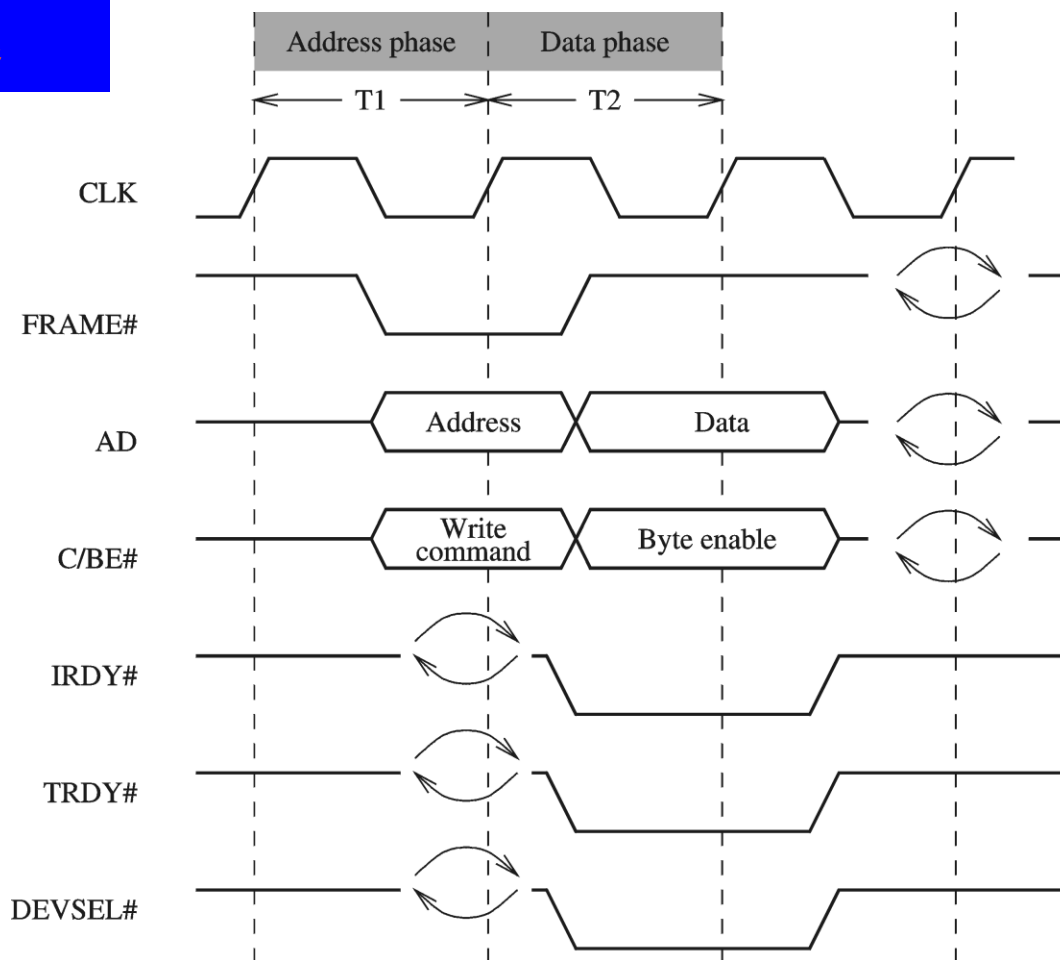
PCI Operations

PCI Read



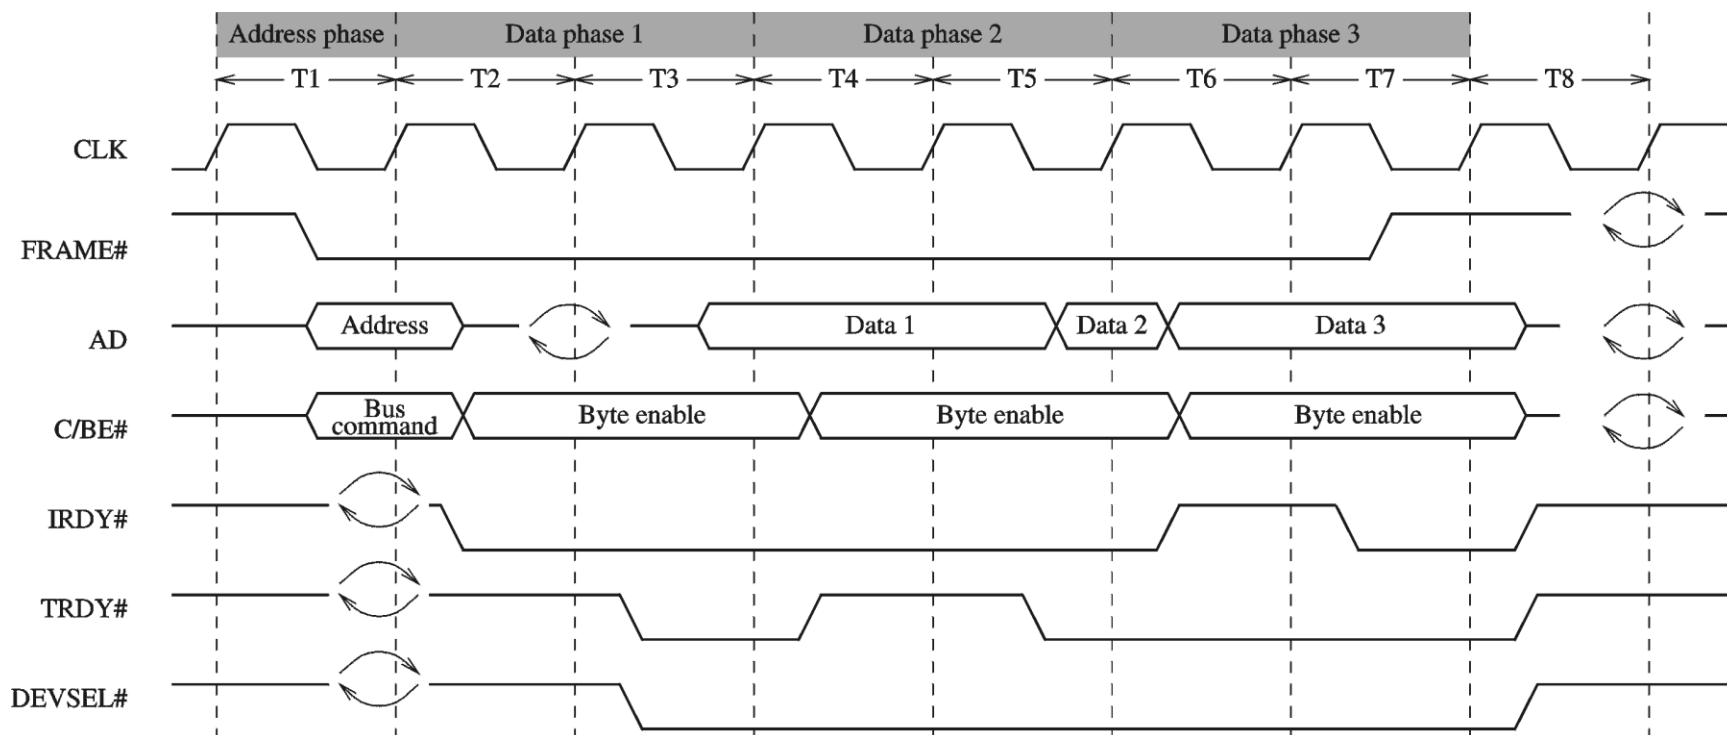
PCI Operations (cont'd)

PCI Write



PCI Operations (cont'd)

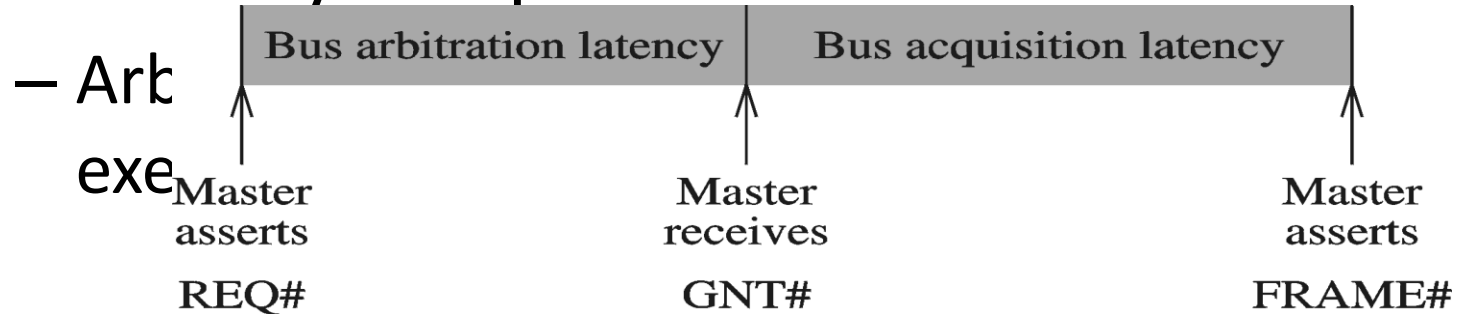
PCI Burst Read Operation



PCI Bus Arbitration

- Uses centralized arbitration
 - Independent grant and request lines
 - REQ# and GNT# lines for each device
 - Does not specify a particular policy
 - Mandates the use of a fair policy

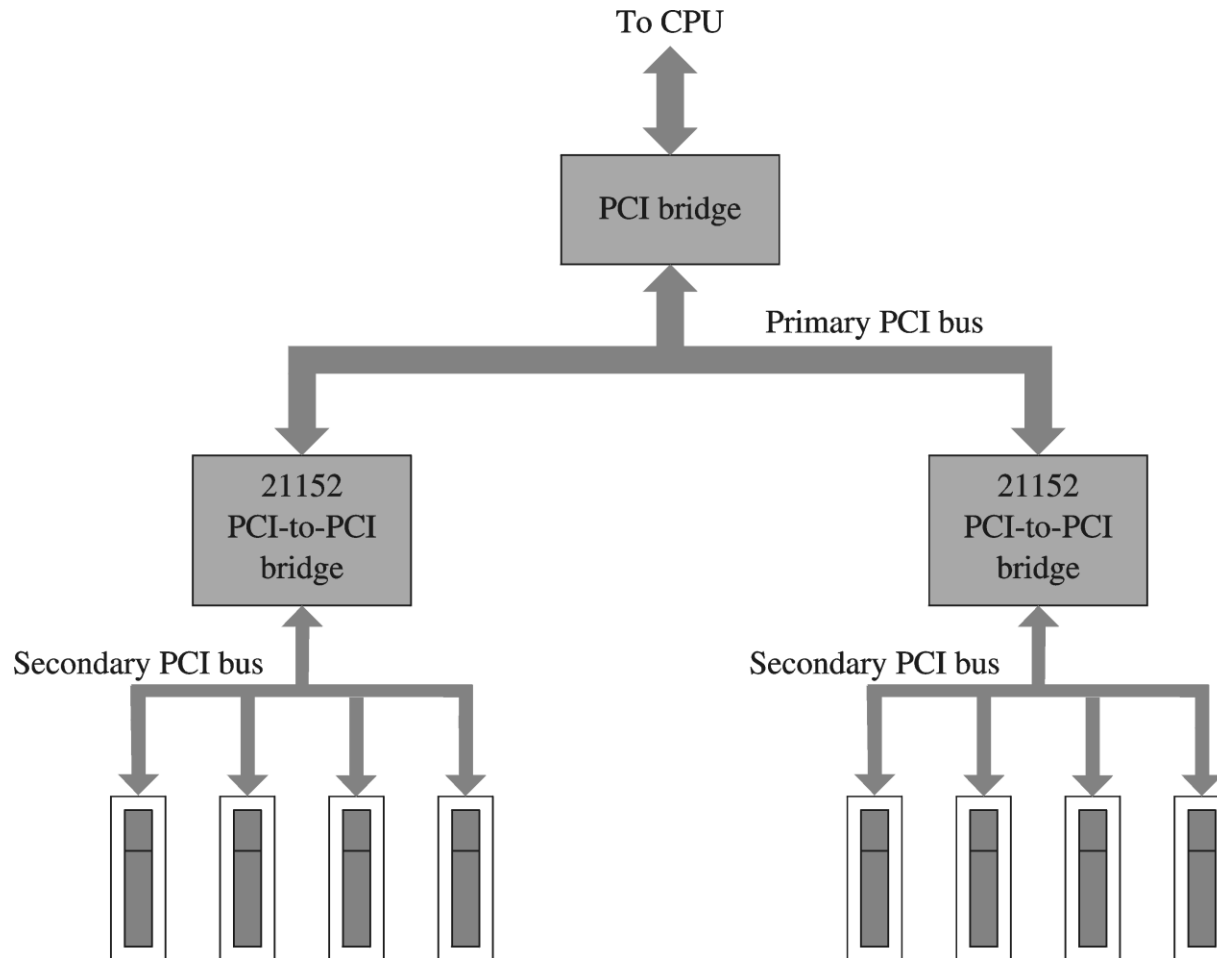
- Two delay components



PCI Bus Hierarchies

- Allows bus hierarchies
- Built using PCI-to-PCI bridges
 - Need two bus arbiters
 - One for the primary bus
 - One for the secondary bus
- Example: Intel 21152 PCI-to-PCI bridge
 - Secondary bus can connect up to 4 devices
 - One internal arbiter available
 - Can be used on the secondary side
 - Need external arbiter for the primary side


PCI Bus Hierarchies (cont'd)




PCI Delays

- 66 MHz implementations pose serious design challenges
- All delays are cut in half compared to 33 MHz clock

33 MHz



66 MHz



AGP

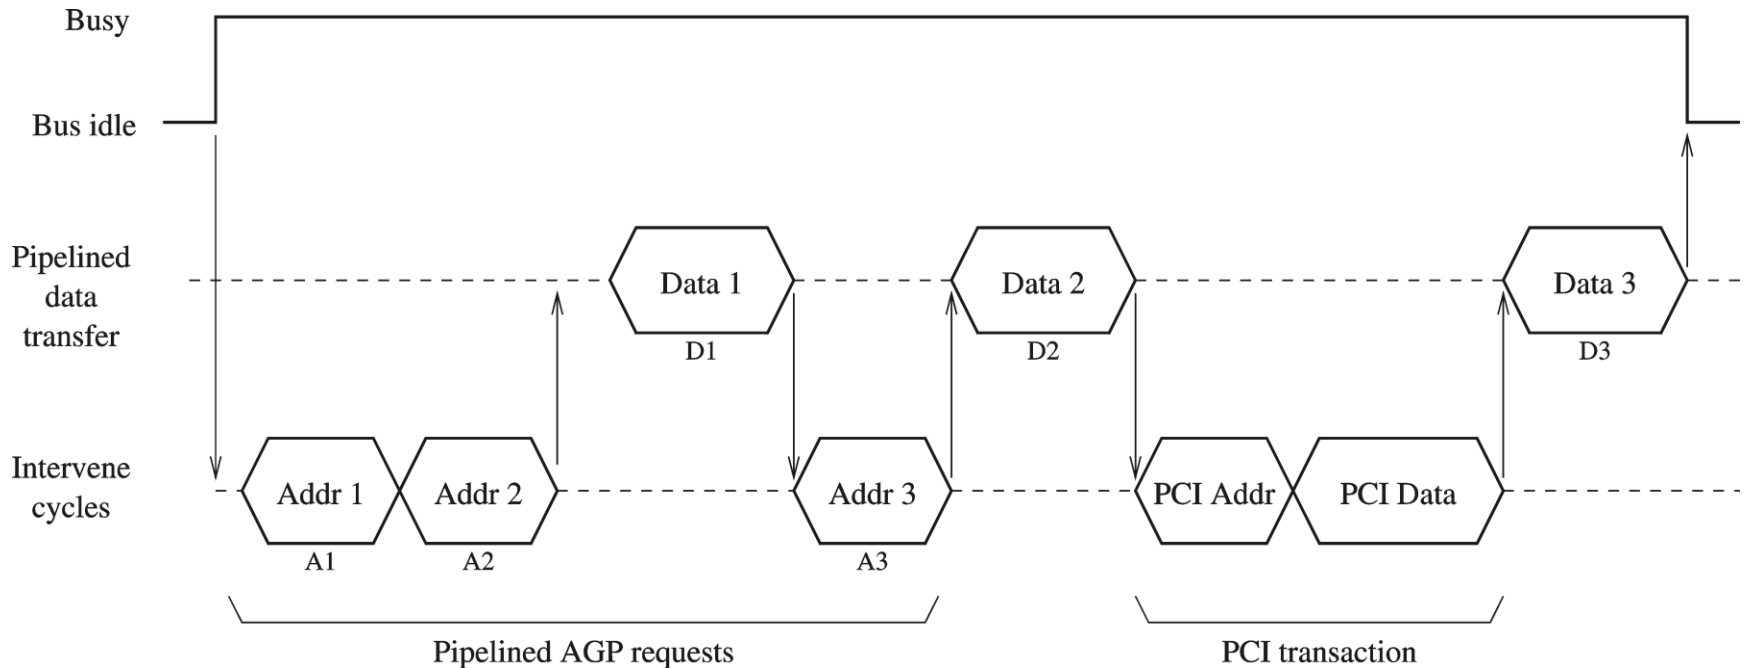
- Supports high-performance video
 - 3D graphics, full-motion video, ...
 - Takes load off the PCI bus
 - Full motion video bandwidth requirement
 - 640 X 480 resolution: 28 MB/s
 - 1024 X 768 resolution: 71 MB/s
 - Twice this if displaying from DVD or hard disk
 - Not a bus, it is a port
 - Connects just two devices
 - CPU and video card

AGP (cont'd)

- AGP specification
 - Based on the 66 MHz PCI 2.1 specification
 - Retains many of the PCI signals
 - 2X mode transmits data on the rising as well as falling edge of clock
 - 4X and 8X speeds are available
- Performance enhancements
 - Uses pipelining
 - Used to hide memory latency
 - Sideband addressing
 - Used to partially demultiplex the bus

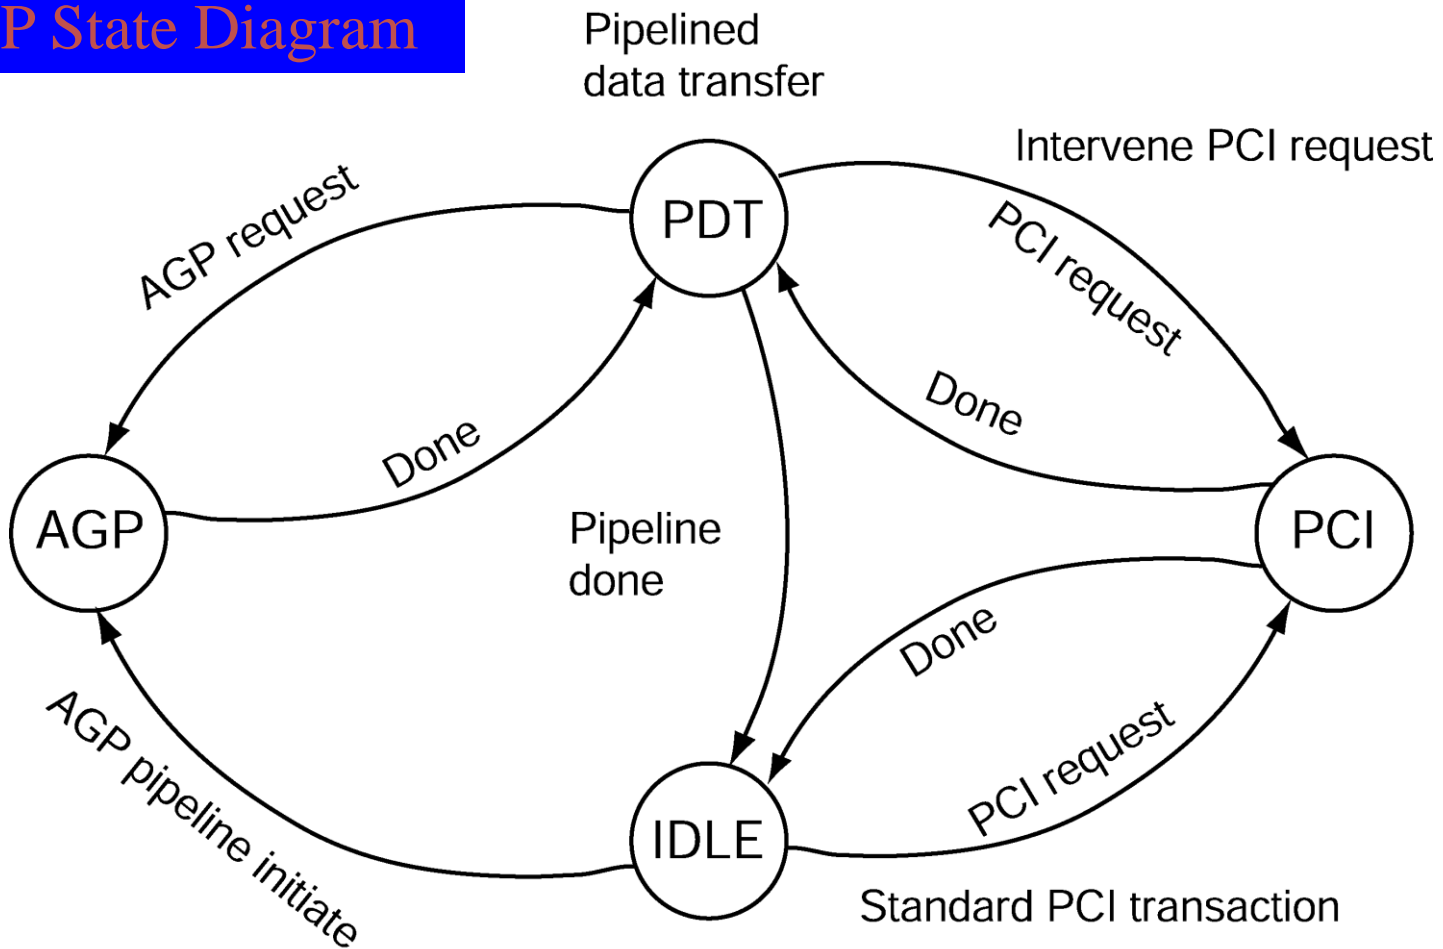
AGP (cont'd)

AGP Pipelined transmission can be interrupted by PCI transactions



AGP (cont'd)

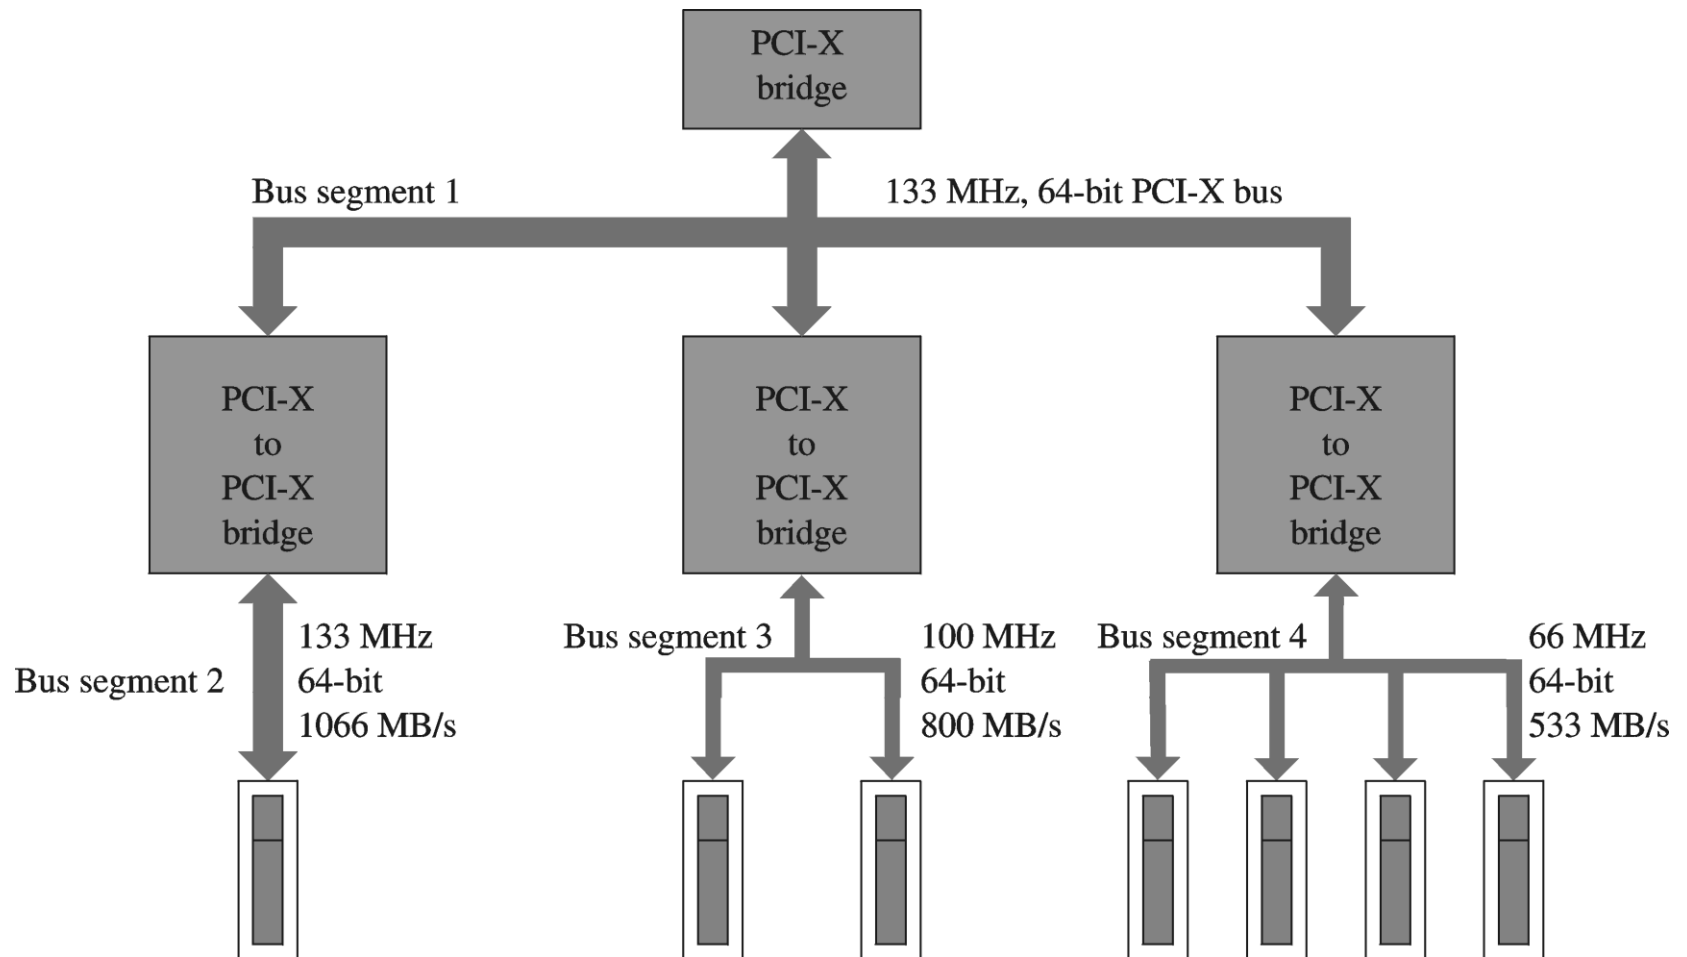
AGP State Diagram



PCI-X Bus

- Addresses the need for higher bandwidth due to
 - Faster I/O buses
 - Faster networks
- Can provide greater than 1 GB/s
 - Achieved by using 64-bit bus operating at 133 MHz
- Uses register-to-register protocol
- Can operate at three different frequencies
 - 66 MHz
 - 100 MHz

PCI-X Bus (cont'd)



PCI-X Bus (cont'd)


- Enhancements include
 - Attribute phase
 - Describes the transaction in more detail than PCI
 - Transaction size, relaxed transaction ordering, identity of transaction initiator
 - Split transaction support
 - PCI: Treats request and reply as a single transaction
 - PCI-X: Splits them into two transactions
 - Optimized wait states
 - Bus can be released due to split transaction support
 - Standard block size movement

PCMCIA Bus


- Started as a standard for memory cards
 - PCMCIA = Personal Computer Memory Card International Association
 - Also called PC Card standard
 - First standard (version 1.0) released in 1990
 - Release 2.0 also supports I/O devices
 - Small form factor
 - 85.5 mm X 54 mm (credit card size)
 - Thickness varies
 - Type I: 3.3 mm thick (used for memory)
 - Type II: 5 mm thick (I/O devices---modems, NICs)
 - Type III: 10.5 mm thick (I/O devices --- hard drives...)

PCMCIA Bus (cont'd)

- PCMCIA connectors (68 pins)
 - Sockets are keyed such that low-voltage card cannot be inserted into a 5 V standard socket



(a) Standard card connector



(b) Low-voltage card connector

PCMCIA Bus (cont'd)

- PCMCIA supports three address spaces
 - Common address space (64 MB)
 - Used for memory expansion
 - Attribute address space (64 MB)
 - Used for automatic configuration
 - I/O address space
 - I/O devices use either Type II or III card
- PC card uses 16-bit data bus
- Also defines a 32-bit standard
 - It is called CardBus

PCMCIA Bus (cont'd)

- Memory Interface

- Address signals

- Address lines (A0 – A25)

- Support 64 MB of address space

- Card Enable Signals (CE1#, CE2#)

- Similar to chip select

- Controls data transfer

- » CE1 = CE2 = High : No data transfer

- » CE1 = low, CE2 = High : Data transfer lower data path (D0 – D7)

- » CE1 = high, CE2 = low : Data transfer upper data path (D8 – D15)

- » CE1 = CE2 = low : 16-bit data transfer (D0 – D15)

PCMCIA Bus (cont'd)

- Memory Interface

- Transaction Signals

- Data lines (D0 – D15)
 - Used to transfer data
 - Output Enable (OE#)
 - Memory read signal
 - Write Enable (WE#)
 - Memory write signal
 - Wait Signal (WAIT#)
 - Can be use to extend the transaction cycle
 - Register Select (REG#)
 - Used to select common memory (high) or attribute memory (low)

Four types of memory transactions:

1. Common memory read
2. Common memory write
3. Attribute memory read
4. Attribute memory write

PCMCIA Bus (cont'd)

– Memory Card Status Signals

- Card Detect Signals (CD1#, CD2#)

CD1#	CD2#	Interpretation
0	0	Card properly inserted
0	1	Card improperly inserted
1	0	Card improperly inserted
1	1	No card inserted

PCMCIA Bus (cont'd)

– Ready/Busy Signal (READY)

- High: indicates the card is ready to be accessed
- Low: Busy executing a command


– Write Protect (WP)

- Gives status of the write protect switch on the card

BVD2 BVD1 Interpretation	
0	0 Battery cannot maintain data integrity
0	1 Battery replacement warning
1	0 Battery cannot maintain data integrity
1	1 Battery is in good condition

PCMCIA Bus (cont'd)

- 16-bit Common memory read cycle
- No wait states



PCMCIA Bus (cont'd)

- I/O Interface
 - Uses Type II or III card
 - Some memory signals are changed for I/O interfacing
 - Some reserved signals are also used for I/O interfacing
 - I/O Read and Write (IORD#, IOWR#)
 - Reserved in memory interface
 - Interrupt Request (IREQ#)
 - Replaces memory READY signal
 - PC card asserts to request interrupt service

PCMCIA Bus (cont'd)

- I/O Size is 16 Bits (IOIS16#)
 - Replaces Write Protect memory signal
 - Low: Indicates I/O is a 16-bit device
 - High: Indicates I/O is a 8-bit device
- System Speaker Signal (SPKR#)
 - Sends audio signal to system speaker
- I/O Status Change (STSCHG#)
 - Replaces BVD1 memory signal
 - Useful in multifunction PC cards
 - Containing memory and I/O functions

PCMCIA Bus (cont'd)

- 16-bit I/O read cycle
- No wait states

